



Parachute: Single-Pass Bi-Directional Information Passing

Mihail Stoian, Andreas Zimmerer, Skander Krid,
Amadou Ngom, Jialin Ding, Tim Kraska, Andreas Kipf

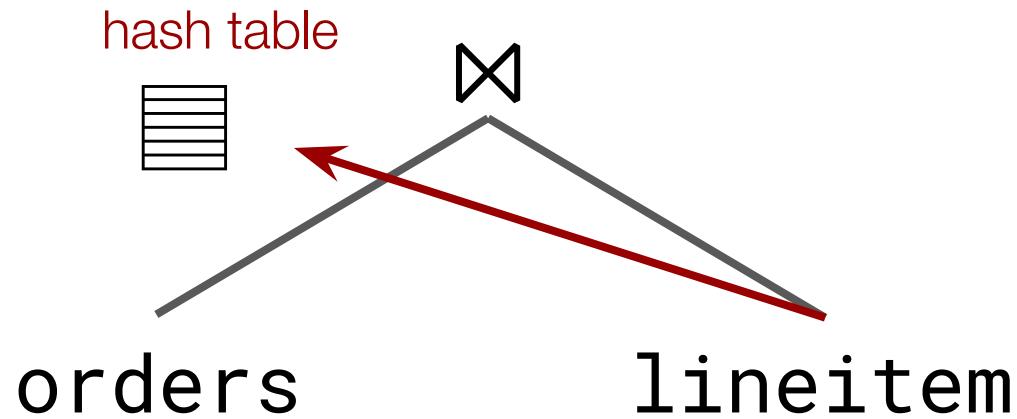
Data Systems Lab @UTN x Data Systems Group @MIT

Databases for OLAP @VLDB'25

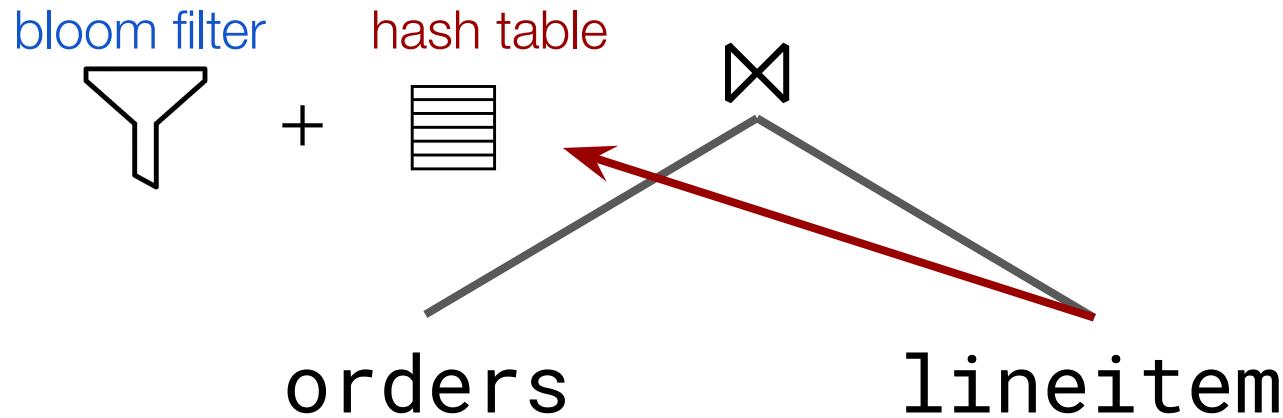
UTN
Technische
Universität
Nürnberg

mit Massachusetts
Institute of
Technology

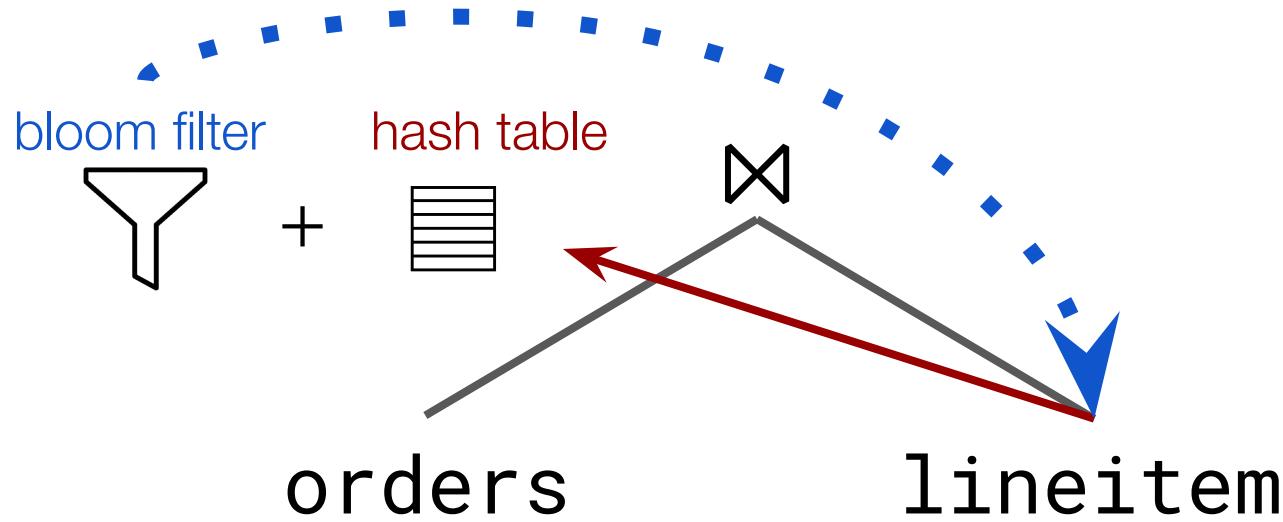
Semi-Join Filtering



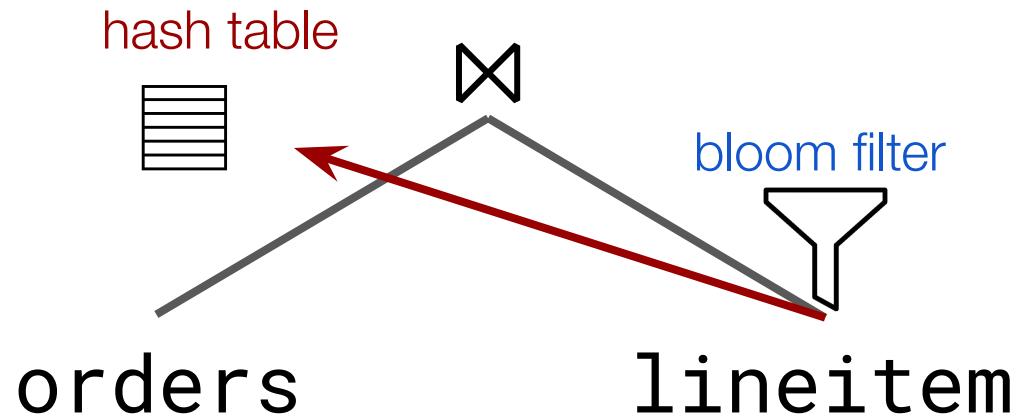
Semi-Join Filtering



Semi-Join Filtering



Semi-Join Filtering



Semi-Join Filtering

- Widely used in production systems, e.g., Amazon Redshift.
- **Goal:** Reduce expensive hash table probes.
- Alternatively: Reduce the #***dangling*** tuples → instance-optimal algorithms.

Semi-Join Filtering

- Widely used in production systems, e.g., Amazon Redshift.
- **Goal:** Reduce expensive hash table probes.
- Alternatively: Reduce the #**dangling** tuples → instance-optimal algorithms.
- Recent research makes instance-optimal algorithms practical:
 - Yannakakis [VLDB'81],
 - Predicate Transfer* [CIDR'24],
 - Lookup & Expand [VLDB'24],
 - Yannakakis++ [SIGMOD'25],
 - Robust Predicate Transfer (RPT) [SIGMOD'25].

Semi-Join Filtering

- Widely used in production systems, e.g., Amazon Redshift.
- **Goal:** Reduce expensive hash table probes.
- Alternatively: Reduce the #**dangling** tuples → instance-optimal algorithms.
- Recent research makes instance-optimal algorithms practical:
 - Yannakakis [VLDB'81],
 - Predicate Transfer* [CIDR'24],
 - Lookup & Expand [VLDB'24],
 - Yannakakis++ [SIGMOD'25],
 - *Robust Predicate Transfer (RPT)* [SIGMOD'25].

So..

- RPT's pull request (<https://github.com/duckdb/duckdb/pull/17326>).



Mytherin commented on May 9 · edited

Collaborator ⋮

Thanks for the PR!

We appreciate the amount of work that has been put into getting this to work - but we need to discuss internally if we adopt this. This is a large change that reworks a lot of the way joins work - and if we adopt this we need to fully understand the strategy and the trade-offs it is making, as well as fully understand the code as we will need to be able to maintain the code going forwards.

As mentioned in the [contributing guidelines](#) - we generally recommend discussing making large changes with the team prior to making them so that we can find the best path forward before a lot of work is done unnecessarily.

We will discuss and let you know.



So..

- RPT's pull request (<https://github.com/duckdb/duckdb/pull/17326>).



Mytherin commented on May 9 · edited

Collaborator ...

Thanks for the PR!

We appreciate the amount of work that has been put into getting this to work - but we need to discuss internally if we adopt this. This is a large change that reworks a lot of the way joins work - and if we adopt this we need to fully understand the strategy and the trade-offs it is making, as well as fully understand the code as we will need to be able to maintain the code going forwards.

As mentor
team prior



We will dis...

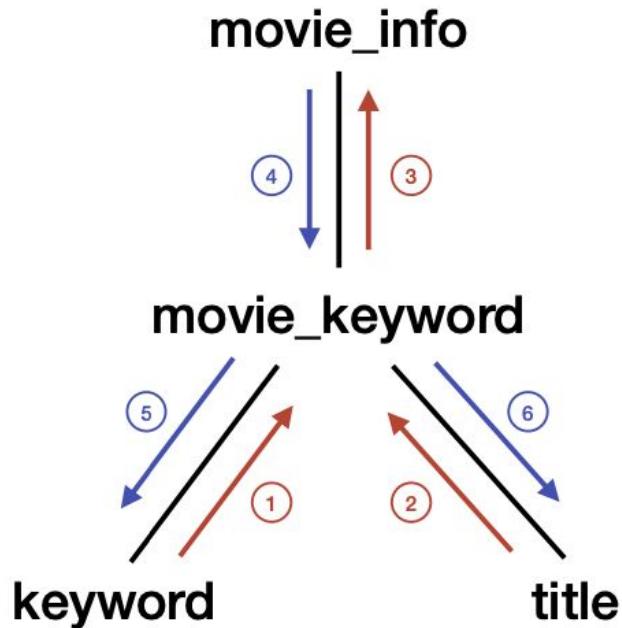


hannes closed this 3 weeks ago



th the
rily.

But Why?



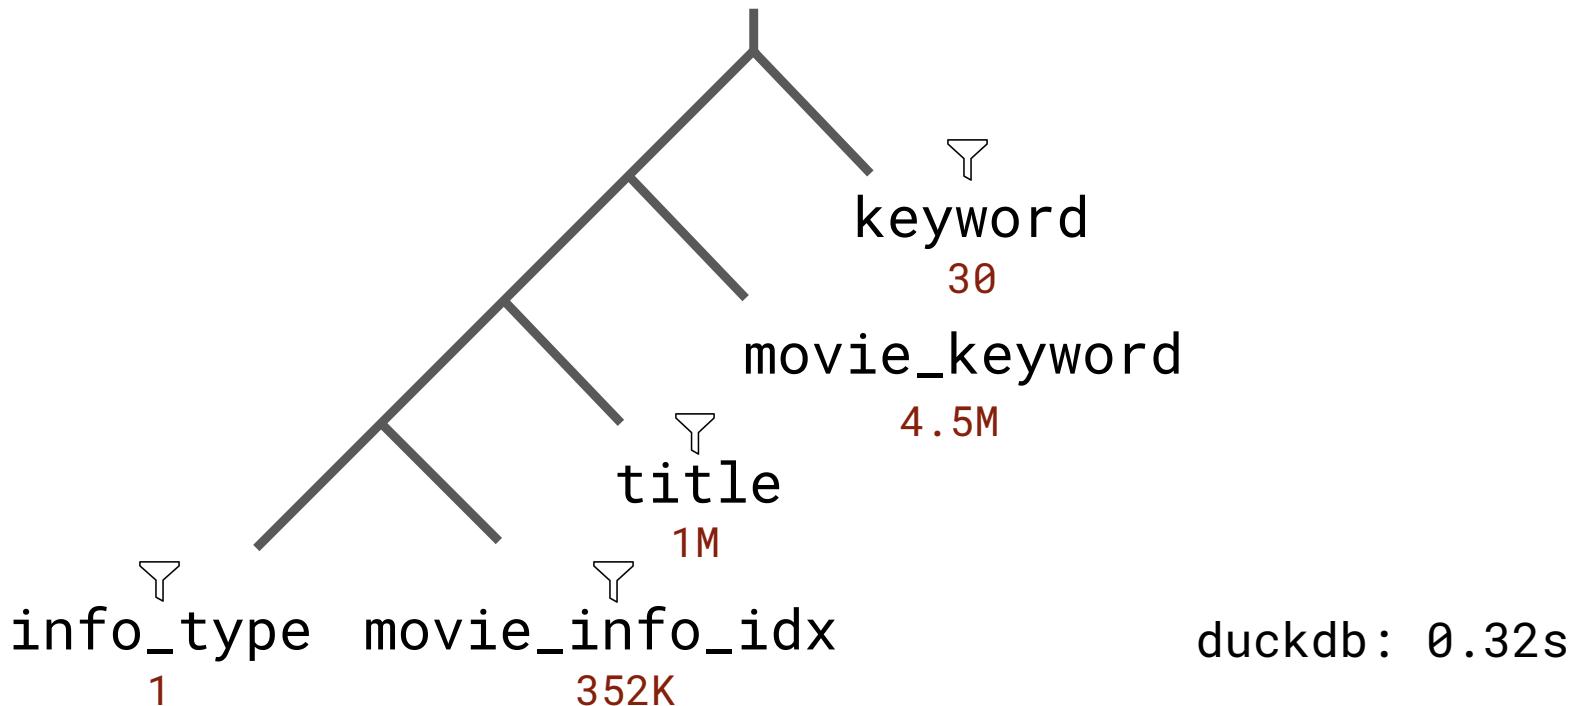
Forward Pass

1. **movie_keyword** ✖ **keyword**
2. **movie_keyword** ✖ **title**
3. **movie_info** ✖ **movie_keyword**

Backward Pass

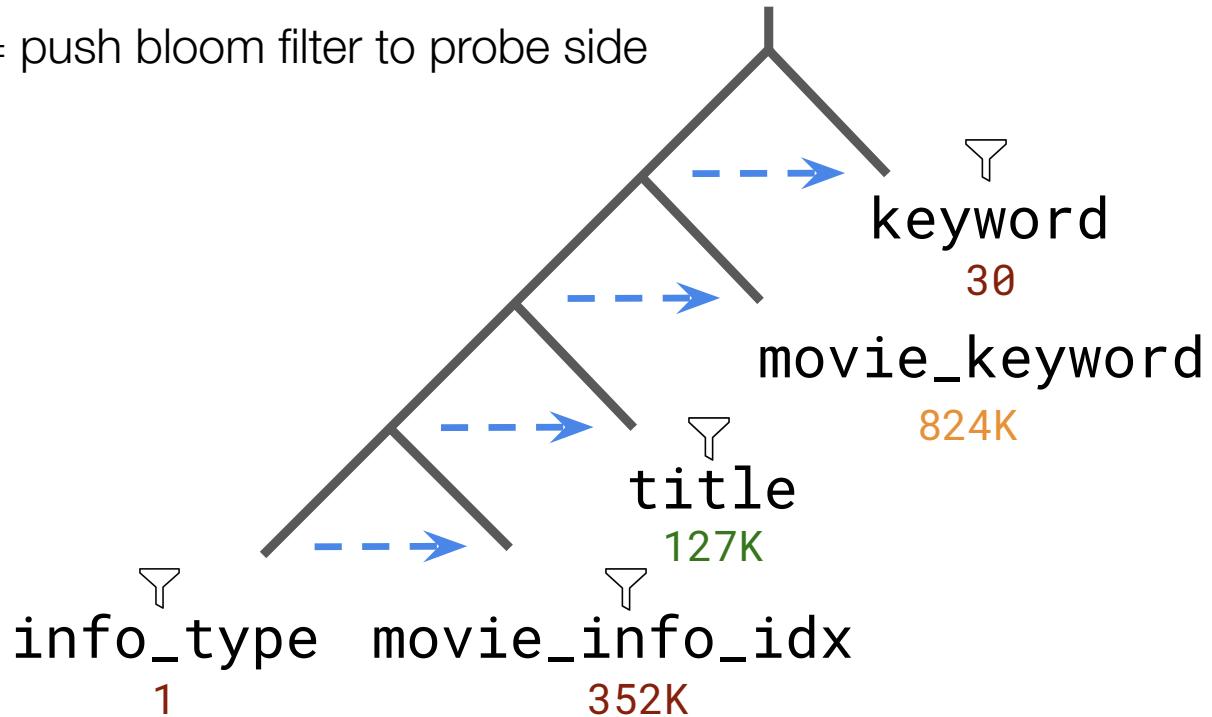
4. **movie_keyword** ✖ **movie_info**
5. **keyword** ✖ **movie_keyword**
6. **title** ✖ **movie_keyword**

Semi-Join Filtering: Why We Love It



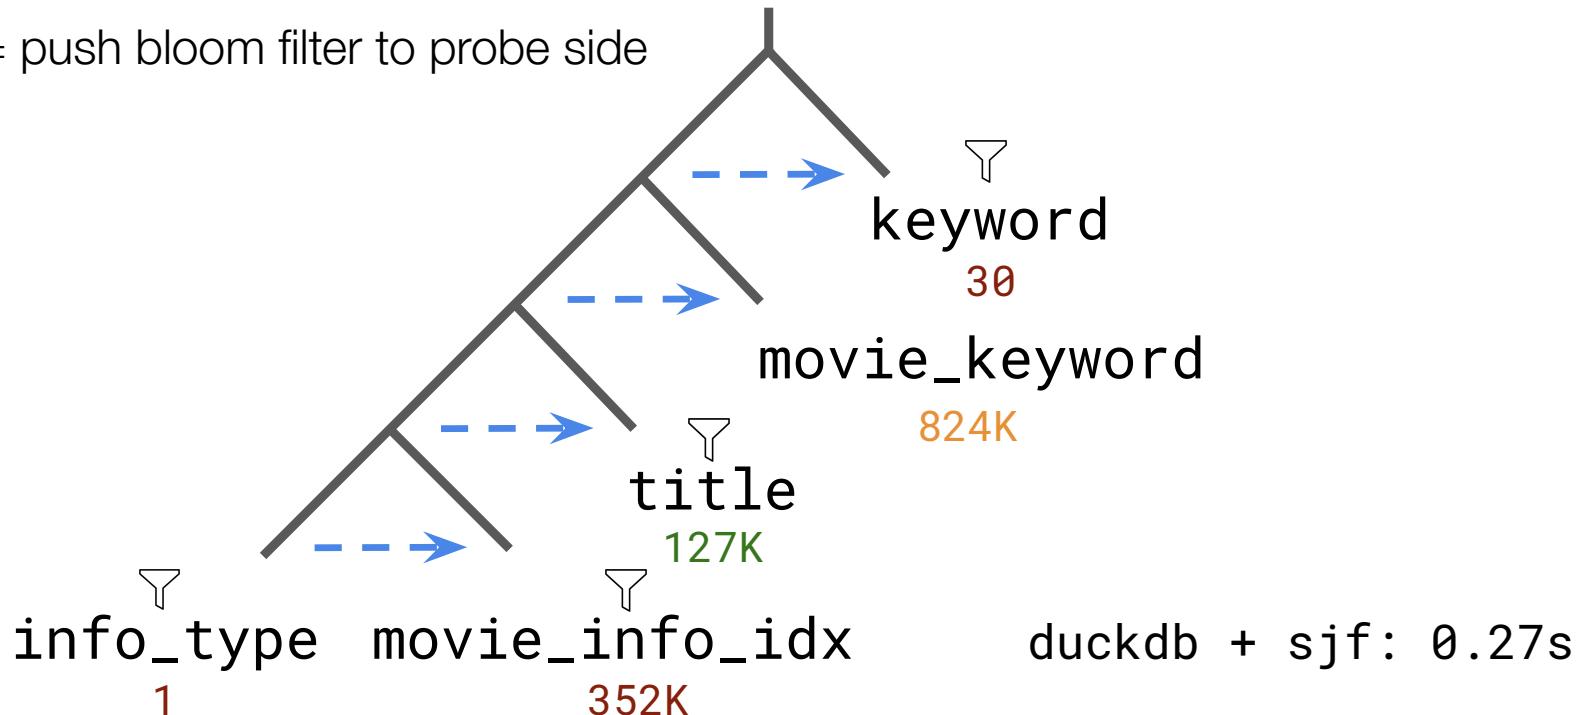
Semi-Join Filtering: Why We Love It

→ = push bloom filter to probe side



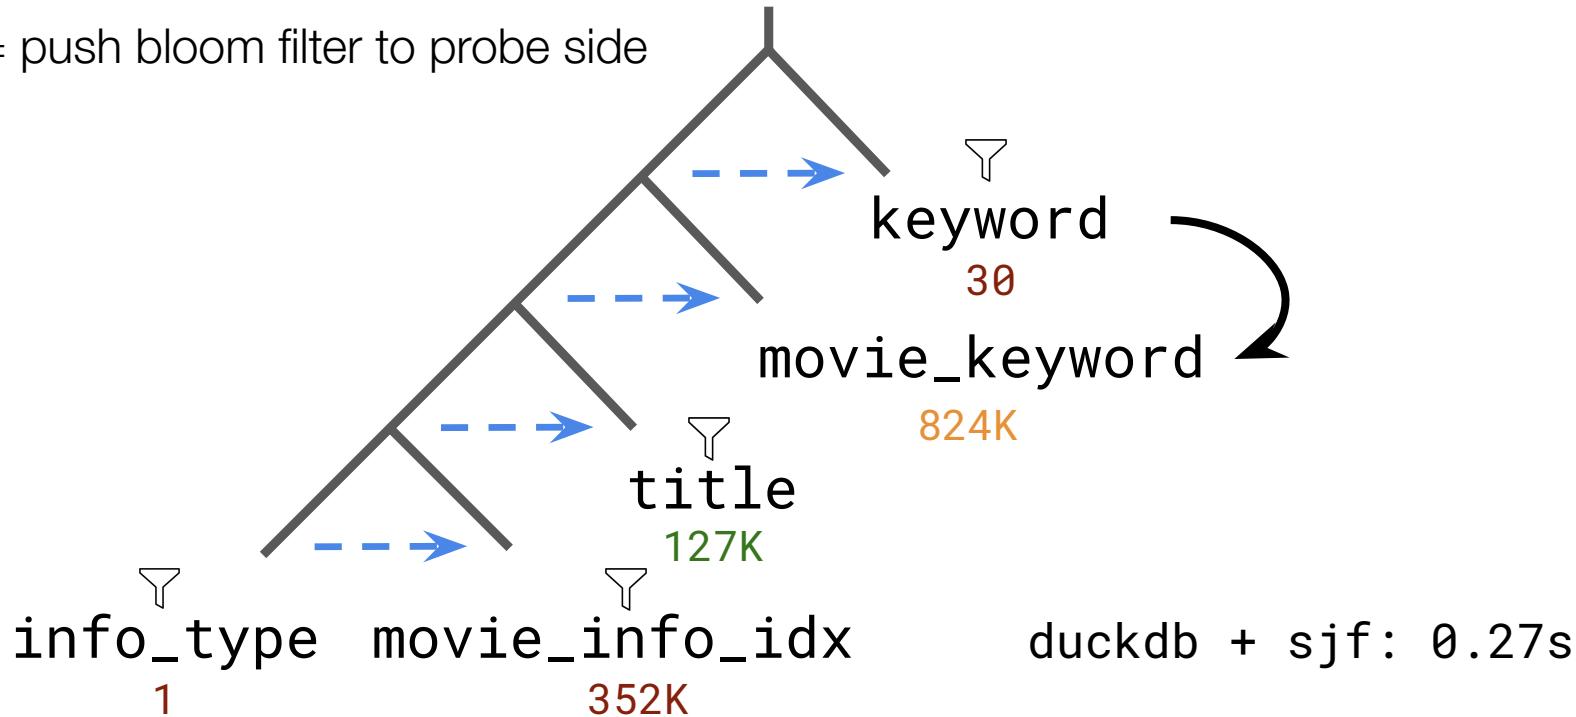
Semi-Join Filtering: Why We Love It

→ = push bloom filter to probe side



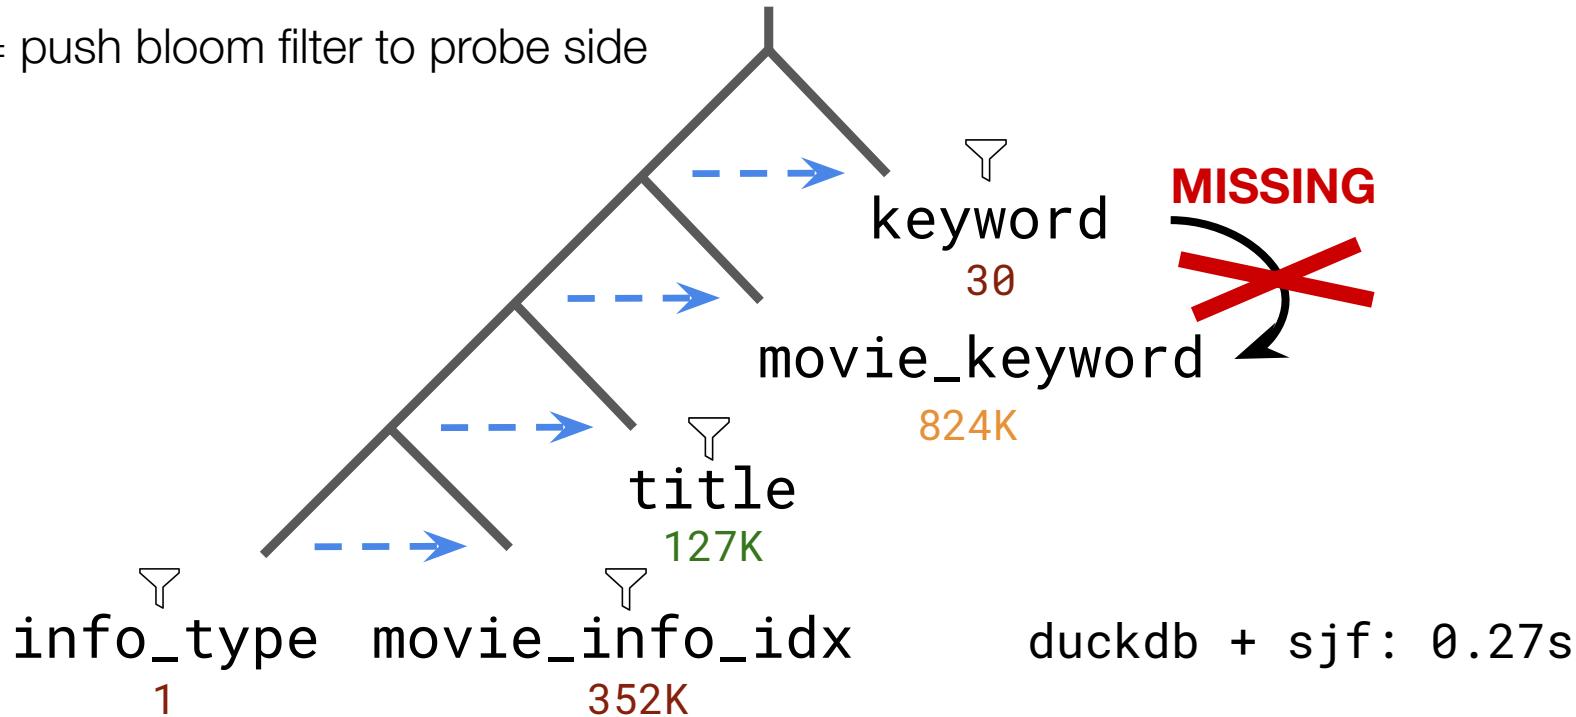
Semi-Join Filtering: *The Missing Bit*

--- → = push bloom filter to probe side



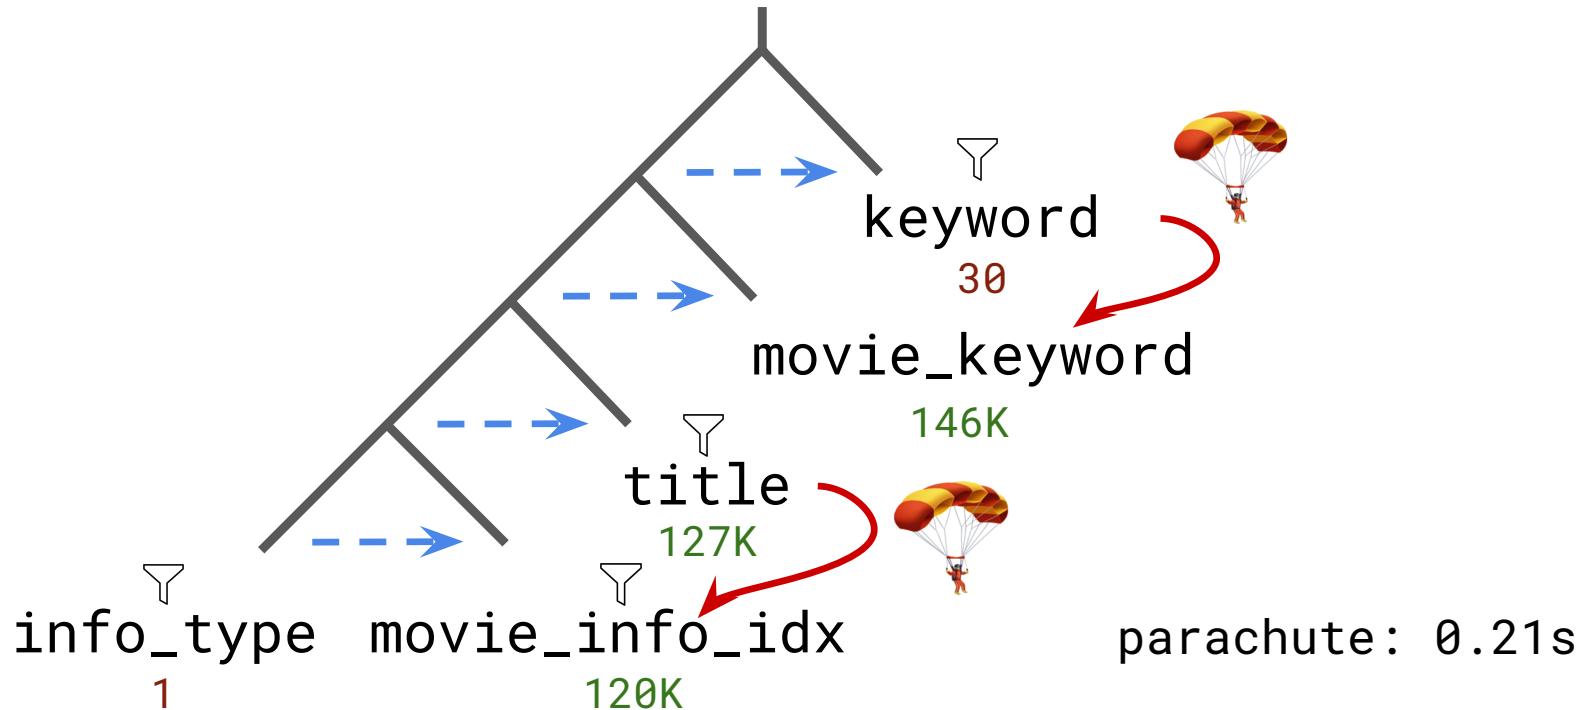
Semi-Join Filtering: *The Missing Bit*

--- → = push bloom filter to probe side





Parachute: *Bi-Directional Information Passing*



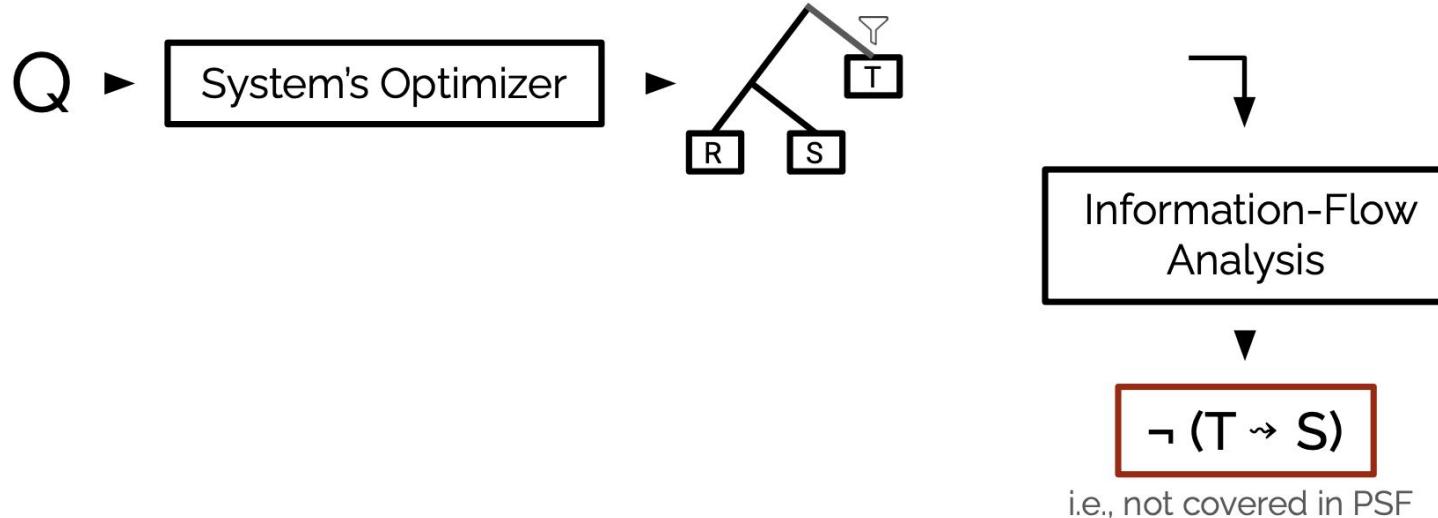


Parachute @Runtime



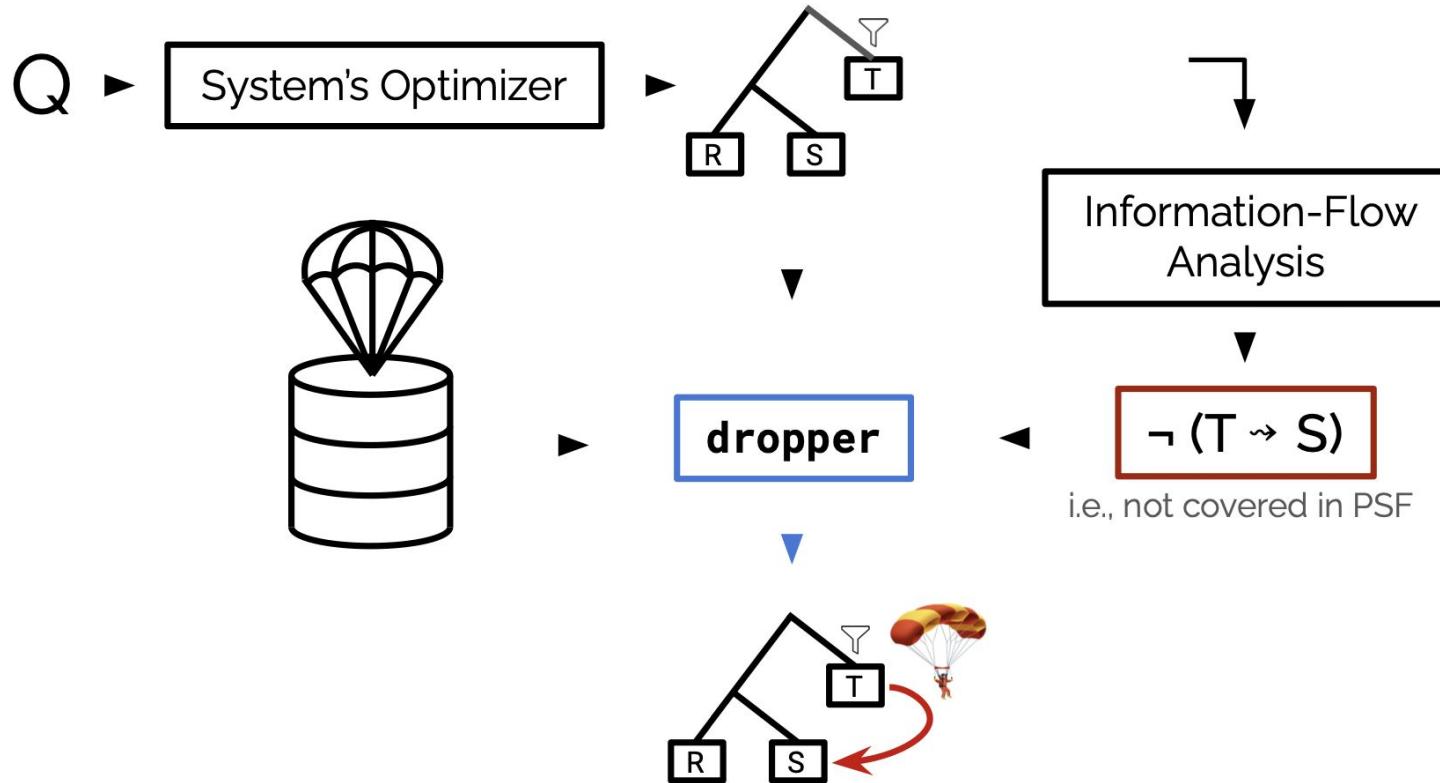


Parachute @Runtime



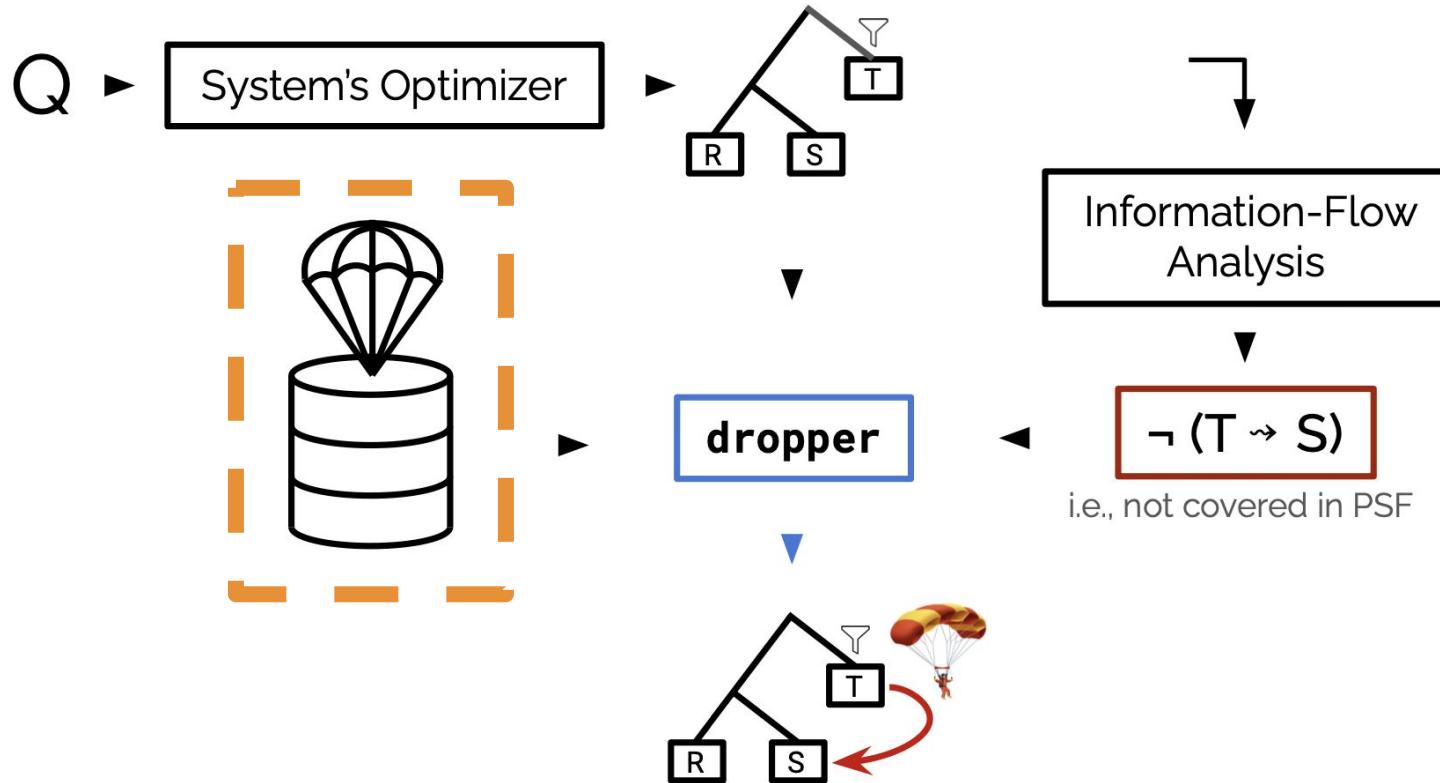


Parachute @Runtime





Parachute @Runtime

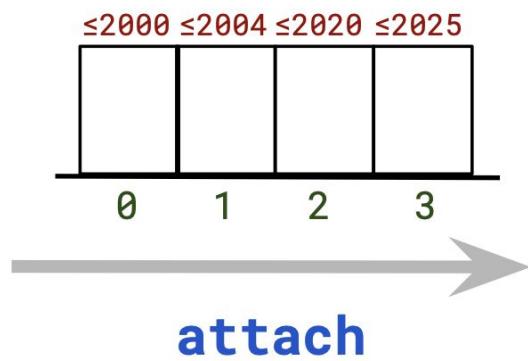




Parachute Columns

id	year
123	1999
456	2018
789	2025

title



movie_id	cast_info
456	1
789	1
789	1

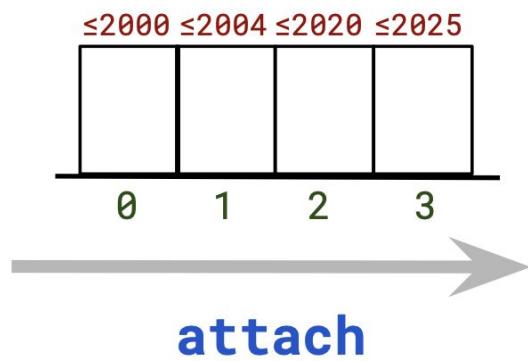
cast_info



Parachute Columns

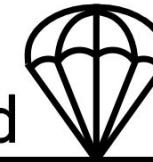
id	year
123	1999
456	2018
789	2025

title



movie_id	
456	1
789	2
789	1

cast_info

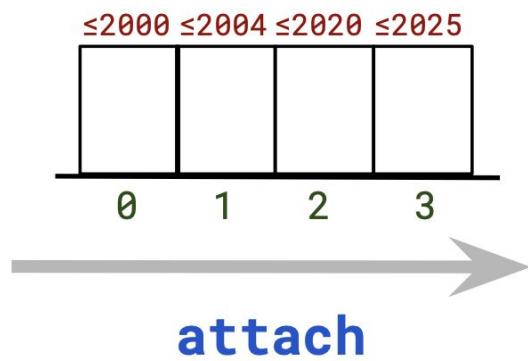




Parachute Columns

id	year
123	1999
456	2018
789	2025

title



movie_id	cast_info
456	1
789	2
789	3

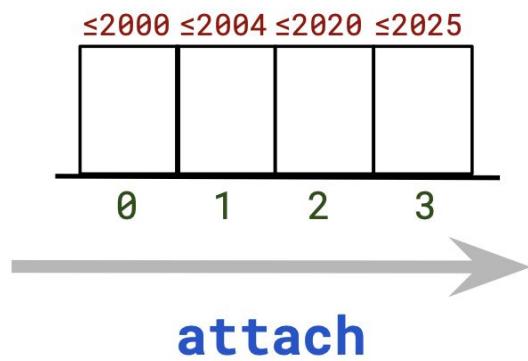
cast_info



Parachute Columns

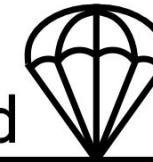
id	year
123	1999
456	2018
789	2025

title



movie_id	
456	1 2
789	3
789	3

cast_info



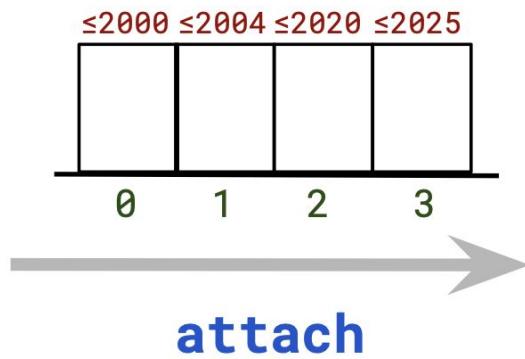


Parachute Predicates

WHERE title.year <= 2016

id	year
123	1999
456	2018
789	2025

title



movie_id	cast_info
456	1 2
789	3
789	3

cast_info

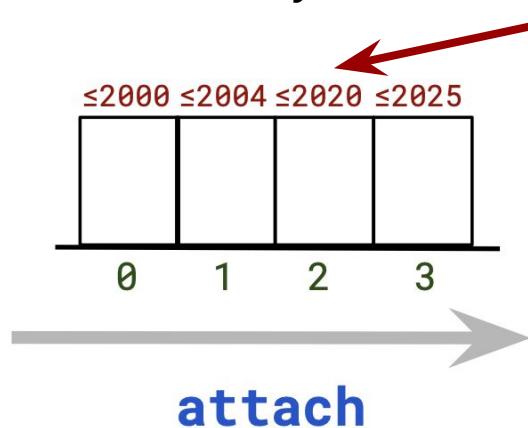


Parachute Predicates

WHERE title.year <= 2016

id	year
123	1999
456	2018
789	2025

title



movie_id	cast_info
456	1 2
789	3 3
789	3

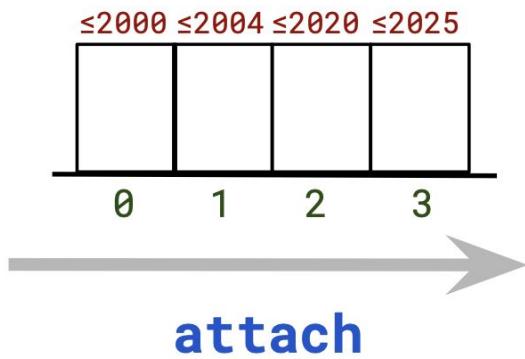
cast_info



Parachute Predicates

WHERE cast_info. <= 2

id	year
title	
123	1999
456	2018
789	2025



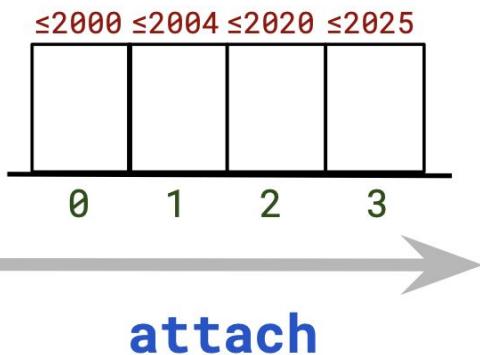
movie_id	cast_info
456	1
789	2
789	3



Parachute Predicates

id	year
title	
123	1999
456	2018
789	2025

WHERE cast_info. <= 2



movie_id		cast_info
456	1	1
789	2	2
789	3	3
789	3	3

Evaluation

- **Benchmarks:**

- JOB: 113 queries.
- CEB: > 13K queries.

Evaluation

- **Benchmarks:**
 - JOB: 113 queries.
 - CEB: > 13K queries.
- **Competitors:**
 - Vanilla DuckDB v1.2: **duckdb v1.2**.
 - DuckDB v1.2 w/ semi-join filtering: **duckdb v1.2 + sjf** (by @andizimmerer).

Evaluation

- **Benchmarks:**
 - JOB: 113 queries.
 - CEB: > 13K queries.
- **Competitors:**
 - Vanilla DuckDB v1.2: **duckdb v1.2**.
 - DuckDB v1.2 w/ semi-join filtering: **duckdb v1.2 + sjf** (by @andizimmerer).
 - Parachute on vanilla DuckDB: **parachute[duckdb]**.
 - Parachute on DuckDB w/ semi-join filtering: **parachute[duckdb + sjf]**.

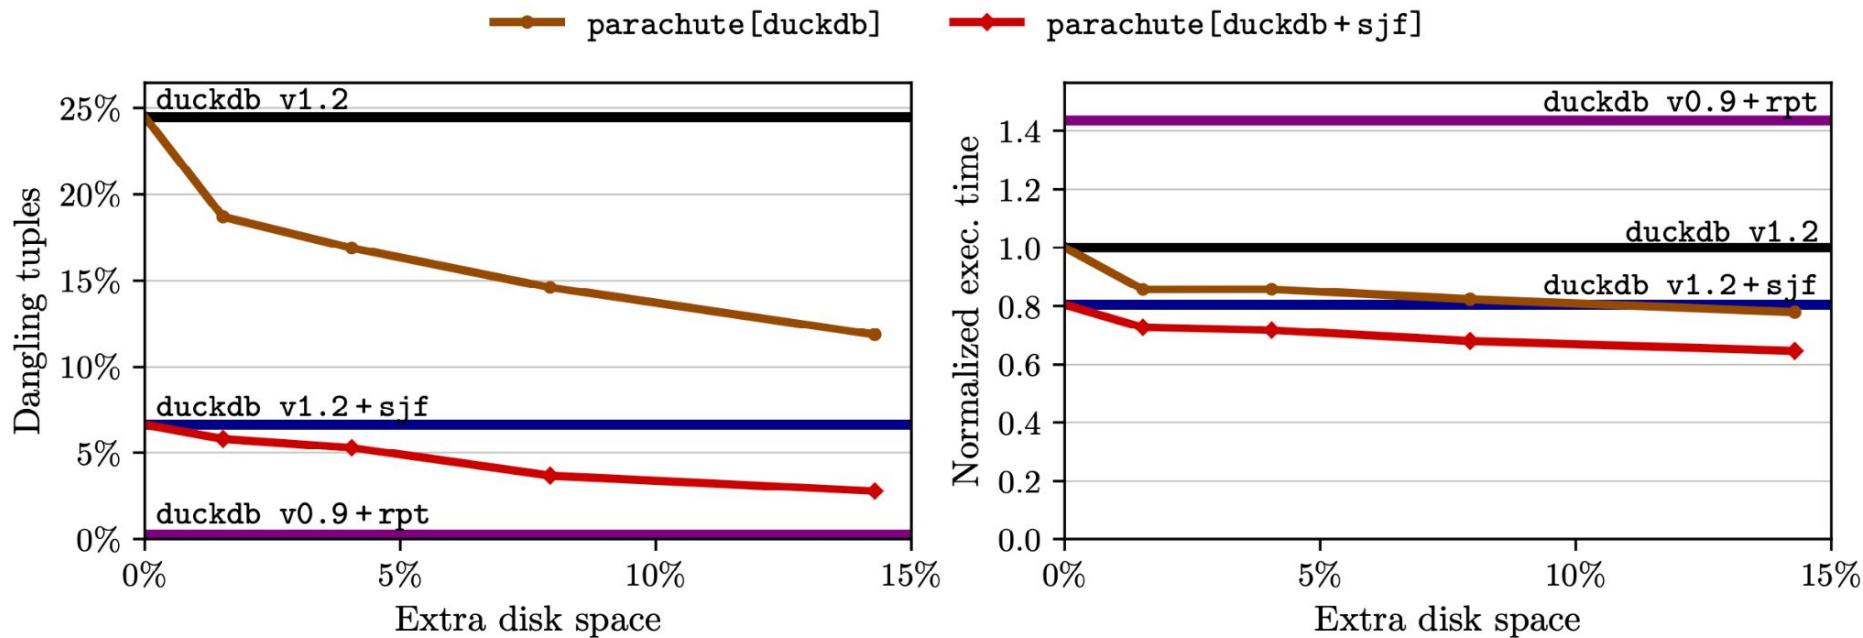
Evaluation

- **Benchmarks:**
 - JOB: 113 queries.
 - CEB: > 13K queries.
- **Competitors:**
 - Vanilla DuckDB v1.2: **duckdb v1.2**.
 - DuckDB v1.2 w/ semi-join filtering: **duckdb v1.2 + sjf** (by @andizimmerer).
 - Parachute on vanilla DuckDB: **parachute[duckdb]**.
 - Parachute on DuckDB w/ semi-join filtering: **parachute[duckdb + sjf]**.
 - *Bonus*: RPT as @SIGMOD'25: **duckdb v0.9 + rpt**.

Evaluation

- **Benchmarks:**
 - JOB: 113 queries.
 - CEB: > 13K queries.
- **Competitors:**
 - Vanilla DuckDB v1.2: **duckdb v1.2**.
 - DuckDB v1.2 w/ semi-join filtering: **duckdb v1.2 + sjf** (by @andizimmerer).
 - Parachute on vanilla DuckDB: **parachute[duckdb]**.
 - Parachute on DuckDB w/ semi-join filtering: **parachute[duckdb + sjf]**.
 - *Bonus*: RPT as @SIGMOD'25: **duckdb v0.9 + rpt**.
- Parachute column bit-widths $\in \{2, 4, 8, 16\}$.

JOB: #DanglingTuples & Exec. Time



Build Overhead

- JOB: Parachute's build amortizes in the 4th run.
- CEB: The overhead is amortized already in the 1st run.

IMDb duckdb's load: 91.67s					
pbw	JOB: 32 parachutes		CEB: 20 parachutes		
	Attach time	Extra space	Attach time	Extra space	
2	244.70s	+1.53%	184.38s	+1.25%	
4	248.44s	+4.03%	187.11s	+3.42%	
8	247.61s	+7.94%	188.47s	+6.55%	
16	265.85s	+14.35%	190.18s	+9.82%	

Future Work

- Parachute columns ~ “cached bloom filters”:
⇒ Parachute columns can be used for partition pruning — think zonemaps.
- Instance-optimized parachute columns.
⇒ Check out *string fingerprints* in AIDB'25 ⇒  1.36x faster LIKE scans.

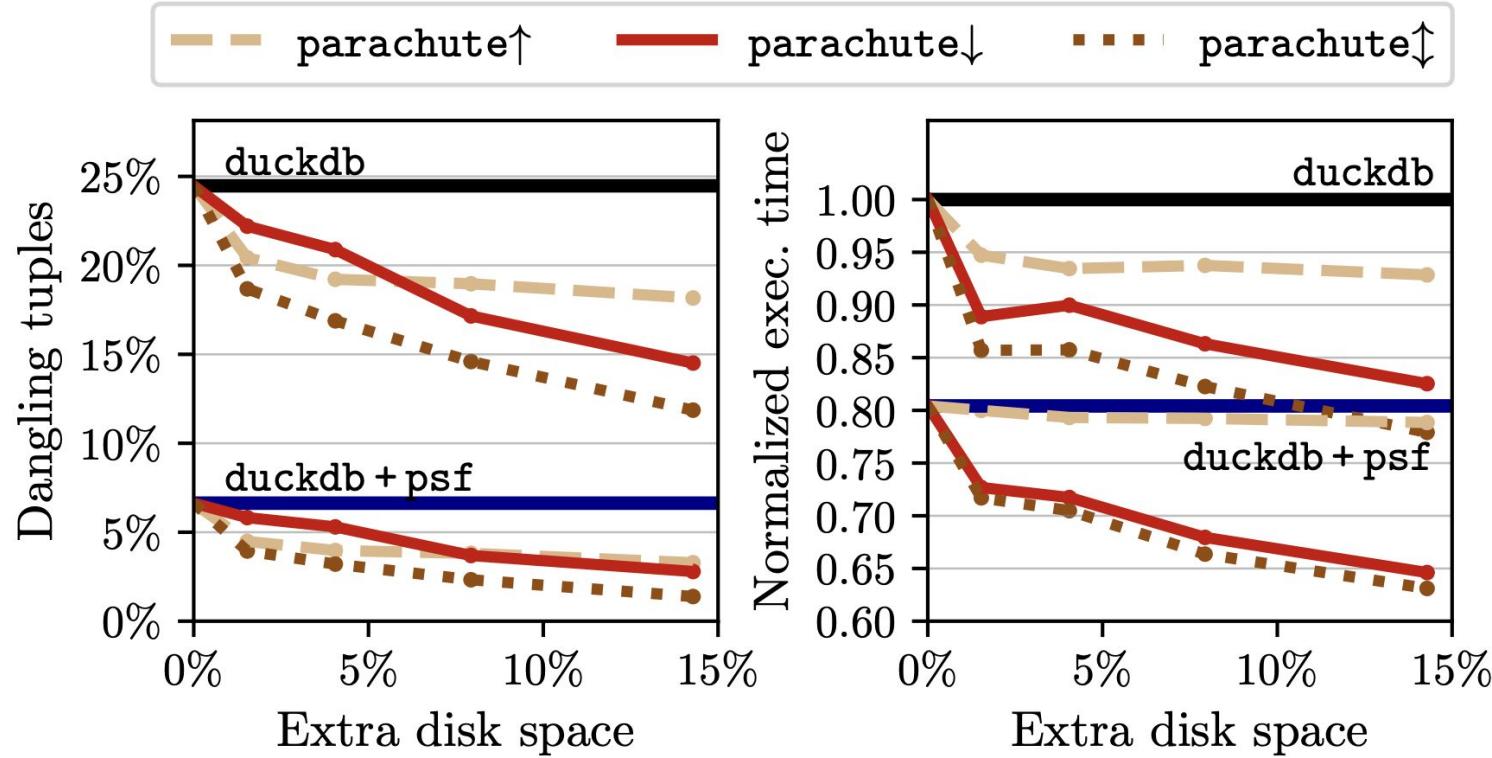


@utndatasystems/parachute

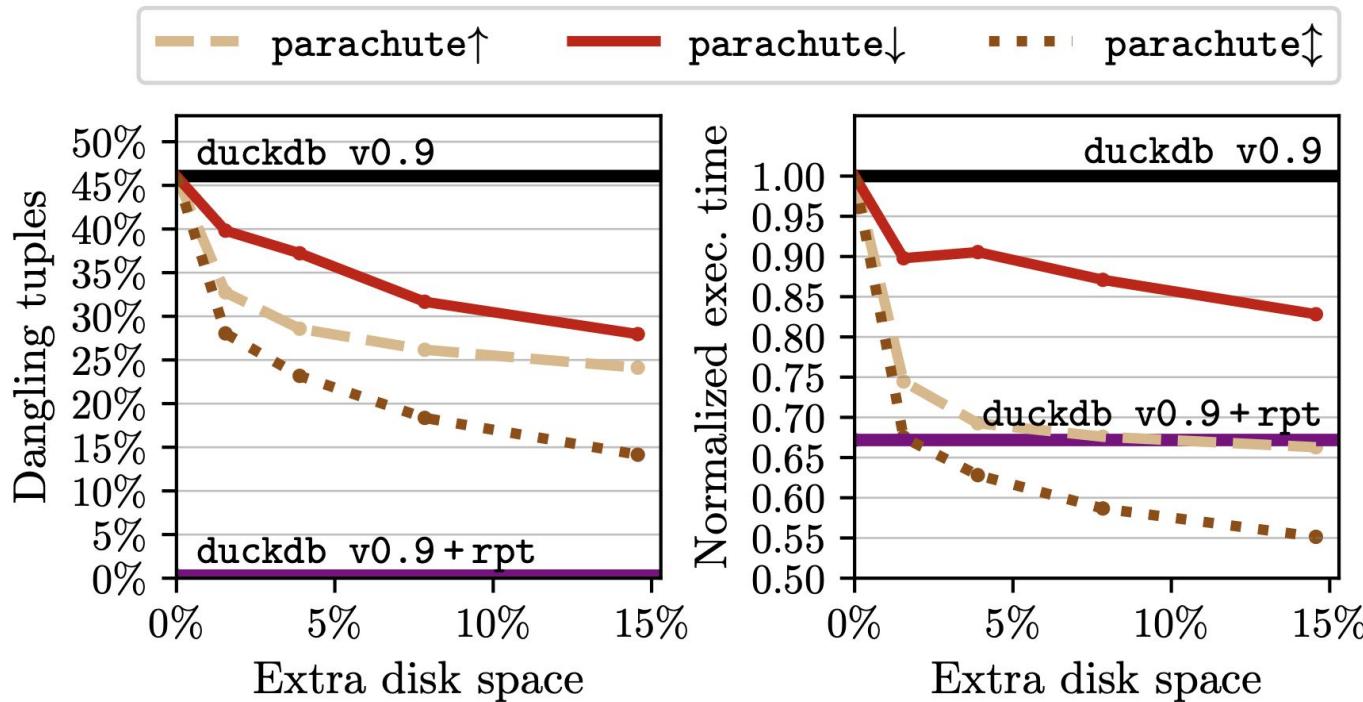
Back-up Slides



Parachute Modes

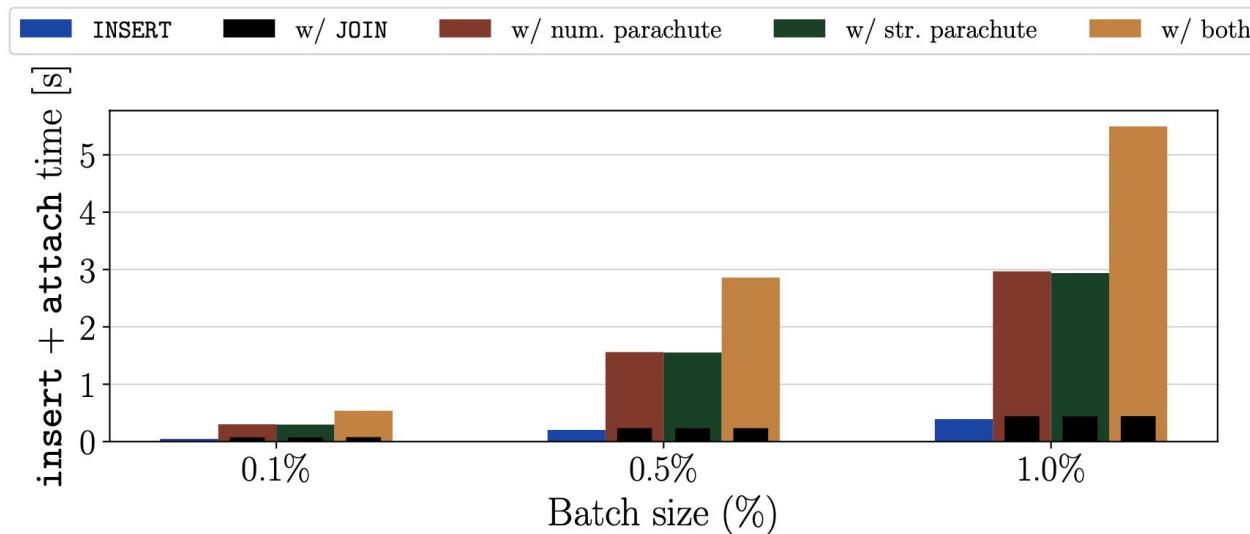


JOB on DuckDB v0.9: Parachute vs. RPT



Update Overhead

- DuckDB v1.2 runs `UPDATE` single-threaded.
⇒ `JOIN` is *not* the main bottleneck.



Information Flow

- **Formalization:**

$$R \rightsquigarrow S \Leftrightarrow (R \prec S) \wedge (R \leftrightarrow S) \wedge \text{is_probe}(S),$$

$$R \rightsquigarrow^{n+1} S \Leftrightarrow \exists T. R \rightsquigarrow^n T \rightsquigarrow S,$$

$$R \rightsquigarrow^* S \Leftrightarrow \exists n > 0. R \rightsquigarrow^n S.$$

- **Precedence Relation:**

$$R \prec S \Leftrightarrow \begin{aligned} & (\text{pipeline}(R) < \text{pipeline}(S)) \\ & \vee (\text{pipeline}(R) = \text{pipeline}(S) \wedge \text{is_probe}(S)). \end{aligned}$$



Parachute's String Fingerprints

1. **Partition** the letter space in a fixed number of bins.
2. Compute a **bitmask** of bin indices

