

MIHAIL STOIAN

Email: mihail.stoian@utn.de
URL: stoianmihail.github.io



WORK EXPERIENCE

Research Assistant <i>University of Technology Nuremberg</i> Data systems lab (Andreas Kipf)	Nov. 2023–present Nuremberg, Germany
Research Intern <i>Gray Systems Lab, Microsoft</i> Robust memory estimation (with Tiemo Bang)	July 2025–Sept. 2025 Barcelona, Spain
Applied Scientist Intern <i>Amazon Redshift</i> Learned systems group (Tim Kraska)	July 2023–Oct. 2023 Munich, Germany
Student Research Assistant <i>TUM, Chair for Database Systems</i> Umbra: A Flash-Based Database System with In-Memory Performance Implementing, improving, and testing the functionality	Mar. 2019–Sept. 2023 Munich, Germany
Student Research Assistant <i>TUM, Chair for Data Analytics and Machine Learning</i> Graph learning with differential privacy	Jan. 2023–Sept. 2023 Munich, Germany
Research Assistant Intern <i>Oracle Labs</i> Graph-in-DB team (Vlad Haprian)	Aug. 2022–Oct. 2022 Zurich, Switzerland
Quantum Software Engineer Intern <i>Infineon Technologies</i> Quantum Algorithms group	Mar. 2021–May 2021 Munich, Germany

EDUCATION

University of Technology Nuremberg <i>PhD, Database Systems</i> Advisor: Andreas Kipf Topic: Robust Query Processing	Nov. 2023–present Nuremberg, Germany
Ludwig-Maximilians-Universität <i>Diploma, Orthodox Theology</i> Pre-diploma (1.9 / 1.0) Topics: Early Church History, Patristics, Liturgies	Oct. 2023–present Munich, Germany
Technical University of Munich <i>M.Sc., Elite Software Engineering</i> Passed with Honors (1.5 / 1.0) Thesis: <i>Optimizing Linearized Dynamic Programming</i> Supervisor: Thomas Neumann	Oct. 2021–Aug. 2023 Munich, Germany
Technical University of Munich <i>M.Sc., Informatics</i> Passed with High Distinction (1.2 / 1.0) Thesis: <i>On the Optimal Linear Contraction Order of Tree Tensor Networks, and Beyond</i>	Oct. 2021–May 2023 Munich, Germany

Supervisor: Christian Mendl

Technical University of Munich

B.Sc., Informatics

Passed with High Distinction (1.2/1.0)

Thesis: *An Efficient Implementation of Polynomial-Time Join Ordering*

Supervisor: Thomas Neumann

Oct. 2018-July 2021

Munich, Germany

AWARDS

SIGMOD Honorable Mention (Best Paper Runner-Up)

2025

DPconv: *Super-Polynomially Faster Join Ordering*

EDBT Best Demonstration Award

2025

Virtual: *Compresing Data Lake Files*

BTW Best Paper Award

2025

Optimizing Linearized Join Enumeration by Adapting to the Query Structure

SIGMOD Student Travel Award

2023

Proposal: "Bridging the Gap Between Computational Fields"

Bronze Medal

2014

National Mathematics Olympiad, Romania

SCHOLARSHIPS

Deutschlandstipendium

Apr. 2022-Mar. 2023

Allianz SE

Munich, Germany

Scholarship awarded by the Ludwig Maximilian University of Munich

RESEARCH PROJECTS

NVIDIA Research

July 2022-Sept. 2023

Student Research Project

Remote

Einsum optimization on GPU

Advisors: Jean Kossaifi, Anima Anandkumar

TUM, Visual Computing & Artificial Intelligence Lab

Apr. 2022-Aug. 2022

Practical Course

Munich, Germany

Outcome: Twofold improvement over DCP, the deep learning approach for iterative closest point (ICP)

Advisor: Matthias Niessner

INTERDISCIPLINARY PROJECTS

INSIGHT

Mar. 2022-Oct. 2022

Chair of Functional Materials (Prof. Peter Müller-Buschbaum)

Munich, Germany

Improved the performance of INSIGHT, the package used by the chair for X-ray measurements

Published in [Journal of Applied Crystallography](#).

PushQuantum

Apr. 2021-Aug. 2021

IQM Quantum Computers

Munich, Germany

[Organiq-Q](#): Quantum simulations for OLED properties ([pitch](#))

PROGRAMMING COMPETITIONS

SIGMOD Programming Contest

Feb. 2022-Apr. 2022

ACM SIGMOD

Munich, Germany

We implemented a blocking system for Entity Resolution

Ranking: 6th place, Team: HyTUM

TECHNICAL SKILLS

Languages: {
 "expert" : { C/C++, Python, SQL, Assembler }
 "advanced" : { Java, Isabelle, HTML/CSS/JS }
}
Frameworks: PyTorch, Spark

LANGUAGE SKILLS

Romanian: Native
English, German: C2
French: C1
Spanish, Ancient Greek: B1

INVITED TALKS

Parachute: Single-Pass Bi-Directional Information Passing. TUMuchData @TUM, January 2026
Instance-Optimized String Fingerprints. Joint Research Seminar, October 2025 (remote)
DPconv: Super-Polynomially Faster Join Ordering. Gray Systems Lab, Microsoft, January 2025 (remote)
Virtual: Compressing World's Parquet Files. TUMuchData @TUM, January 2025
What Selinger Forgot to Tell You About Query Optimization. Systems Group, TU Darmstadt, June 2024 (remote)
What do databases and tensor networks have in common? University of Jena, August 2023

PREPRINTS

[xBound: Join Size Lower Bounds](#)
Mihail Stoian, Tiemo Bang, Hangdong Zhao, Jesús Camacho-Rodríguez, Yuanyuan Tian, Andreas Kipf
tba

[Redbench: Workload Synthesis From Cloud Traces](#)
Johannes Wehrstein, Roman Heinrich, **Mihail Stoian**, Skander Krid, Martin Stemmer, Andreas Kipf, Carsten Binnig, Muhammad El-Hindi
tba

PUBLICATIONS

[Mind the Gap. Doubling Constant Parametrization of Weighted Problems: TSP, Max-Cut, and More](#)
Mihail Stoian
STACS 2026

[Waiting to Decompress: The Economics of LLM-Based Compression](#)
Andreas Kipf, Tobias Schmidt, Ping-Lin Kuo, Skander Krid, Moritz Rengert, Luca Heller, Andreas Zimmerer, **Mihail Stoian**, Varun Pandey, Alexander van Renen
CIDR 2026

[Parachute: Single-Pass Bi-Directional Information Passing](#)
Mihail Stoian, Andreas Zimmerer, Skander Krid, Amadou Latyr Ngom, Jialin Ding, Tim Kraska, Andreas Kipf
VLDB 2025

[Instance-Optimized String Fingerprints](#)
Mihail Stoian*, Johannes Thürauf*, Andreas Zimmerer, Alexander van Renen, Andreas Kipf
AIDB @VLDB 2025

[Redbench: A Benchmark Reflecting Real Workloads](#)
Skander Krid, **Mihail Stoian**, Andreas Kipf
aiDM @SIGMOD 2025

[DPconv: Super-Polynomially Faster Join Ordering](#)
Mihail Stoian, Andreas Kipf
SIGMOD 2025

[Virtual: Compressing Data Lake Files](#)

Mihail Stoian, Alexander van Renen, Jan Kobiolka, Ping-Lin Kuo, Andreas Zimmerer, Josif Grabocka, Andreas Kipf
EDBT 2025

[Optimizing Linearized Join Enumeration by Adapting to the Query Structure](#)

Altan Birler, **Mihail Stoian**, Thomas Neumann
BTW 2025

[Lightweight Correlation-Aware Table Compression](#)

Mihail Stoian, Alexander van Renen, Jan Kobiolka, Ping-Lin Kuo, Josif Grabocka, Andreas Kipf
3rd Table Representation Learning Workshop (TRL@NeurIPS), 2024

[Unified Mechanism-Specific Amplification by Subsampling and Group Privacy Amplification](#)

Jan Schuchardt, **Mihail Stoian***, Arthur Kosmala*, Stephan Günemann
37th Conference on Neural Information Processing Systems (NeurIPS), 2024

[On the Optimal Contraction Order of Tree Tensor Networks, and Beyond](#)

Mihail Stoian, Richard Milbradt, Christian B. Mendl
SIAM Journal on Scientific Computing, 2024

[Approximate Min-Sum Subset Convolution](#)

Mihail Stoian

22nd International Workshop on Approximation and Online Algorithms (WAOA), 2024

[DataLoom: Simplifying Data Loading with LLMs](#)

Alexander van Renen, **Mihail Stoian**, Andreas Kipf
Proceedings of the VLDB Endowment, Vol. 17, 2024

[Corra: Correlation-Aware Column Compression](#)

Hanwen Liu, **Mihail Stoian**, Alexander van Renen, Andreas Kipf
2nd Workshop on Cloud Databases (CloudDB @VLDB), 2024

[Fast Joint Shapley Values](#)

Mihail Stoian

Student Research Competition, Companion of the International Conference on Management of Data, 2023

[Faster FFT-based Wildcard Pattern Matching](#)

Mihail Stoian

Student Research Competition, Companion of the International Conference on Management of Data, 2023

[Concurrent Link-Cut Trees](#)

Mihail Stoian

Student Research Competition, Proceedings of the International Conference on Management of Data, 2022

[PLEX: Towards Practical Learned Indexing](#)

Mihail Stoian, Andreas Kipf, Ryan Marcus, Tim Kraska
3rd International Workshop on Applied AI for Database Systems and Applications (AIDB), 2021

[Benchmarking Learned Indexes](#)

Andreas Kipf, Ryan Marcus, Alexander van Renen, **Mihail Stoian**, Sanchit Misra, Alfons Kemper, Thomas Neumann, Tim Kraska
Proceedings of the VLDB Endowment, Volume 14, 2021

[RadixSpline: A Single-Pass Learned Index](#)

Andreas Kipf, Ryan Marcus, Alexander van Renen, **Mihail Stoian**, Alfons Kemper, Tim Kraska, Thomas Neumann
3rd International Workshop on Exploiting AI Techniques for Data Management (aiDM), 2020

[SOSD: A Benchmark for Learned Indexes](#)

Andreas Kipf, Ryan Marcus, Alexander van Renen, **Mihail Stoian**, Alfons Kemper, Tim Kraska, Thomas Neumann
NeurIPS Workshop on Machine Learning for Systems, 2019